

Tab 1

Here are **credible business models** for **Zero (iOS)** built like a Harvard Business School MBA advising across multiple startup patterns.

1.

Business Model Archetypes for Zero

A. Subscription-First SaaS (Core Revenue)

- Primary Rev-stream: \$9.99/mo or \$99.99/yr
- Includes:
 1. Mail vs Ads classification
 2. Thread + email summaries
 3. 10 core 1-tap actions
 4. Offline + resilience features
 5. High-quality AI tuning loop access (non-destructive, parent-tuned models)

Why it works:

- Parents will pay for peace of mind + time back
 - Annual pricing encourages retention
 - Margin improves with scale + batching cloud costs
 - Feature unlock progression can be gated by subscription tier
-

B. Action Pack Add-Ons (Modular Expansion Revenue)

Sell action bundles on top of subscription:

Pack	Price	Contents
School	+\$2.99/mo	forms, reminders, calendar, signatures
Shopping	+\$3.99/mo	price tracking, recurring purchase scheduling
Travel	+\$2.99/mo	flight check-in, status summaries, wallet passes
Bills	+\$4.99/mo	invoice pay, autofill, summarization

Benefits:

- ARPU increases without cluttering core UX
- Users only pay for what they activate (optionality)
- Compound actions justify higher add-on pricing in future sprints

C. Enterprise White-Label Family Plan (B2B2C)

License to existing parenting ecosystems:

- **Schools and districts:** Zero manages digital consent & classified parent outreach
- **Airlines loyalty:** One-tap check-ins integrated, co-marketed
- **Package shippers:** Embedding tracking intelligence

- **Credit card networks:** Integrate bill-pay actions adjacent to card approvals
- **Major retailers:** Sale + recurring purchases + sponsored tunable actions

Pricing example: \$4/user/mo to partners, 12-mo contract minimum.

D. Sponsored Actions + ML Feedback Economy (Non-Primary but Huge Upside)

- Brands pay **per completed action swipe** rather than impressions
- “Tuned suggestions” appear **when confidence ≥4**
- No action bypasses parent override
- CPC-like pricing but **CPA based on completion**

Example:

- \$2–\$7 per completed sponsored action
- Only 1 action suggestion shown at a time, from curated brand registry
- We learn while users swipe → ML tuning data pipeline

This makes marketing item 1-5 *demonstrably part of value* without feeling like ads.

E. Hybrid Tokenized Credibility (Analytics Only, Not Crypto UX For App Yet)

You later incorporate a reputation/credibility score for summaries + action quality based on:

- parent edits
- conflict overrides

- acceptance rates
- failure logs
- engagement with model tuning

This becomes a non-user-facing model versioning layer that could inform pricing elasticity, churn risk, and quality scores.

2.

Revenue Sizing & Cohort Burn Alignment

Model	ARPU	Notes
Sub + 1 add-on	\$13.97	typical busy parent ARPU
Sub + 3 add-ons	\$20.97	power parent ARPU
Enterprise White-Label	\$4/use r	12-mo contracts, stable margin

Projected path at scale:

- 10k parents → \$2.5M ARR
- 50k parents → \$12.5M ARR
- 200k parents → \$50M ARR
- 1M parents → \$250M+ ARR

Assumes 70% annual retention after PMF, 90%+ classifier, 98% action reliability.

3.

Cost → Revenue Coupling With Feature 1-5

Cost centers injected into product economics:

- AI model calls (classification + summaries + actions)
- App-engineering payroll
- QA + Accessibility + Ops
- Marketing install pushes
- White-label licensing infra
- Feedback quality training + observability tooling

We monitor cost per interaction via built-in cost telemetry.

4. 1-Week Sprint ✓ Story Acceptance To Land 1-5

Example grooming for action reliability sprint (~2 weeks):

- 10 actions (10 SP each at avg)
- summary accuracy guardrails (8 SP)

- classifier p95 improvements (5 SP)
- fallback UX coverage (5 SP)
- cost telemetry wiring (3 SP)

Total committed per sprint: 29-36 SP with 20% buffer.

5. Recommended Artifacts (Next To Build)

Artifact	Audience	Purpose
12-slide + speaker notes deck	Exec/Investors	Vision + demos
Notion epics + 10 actions	Internal team	Traceability
Burn model sheet	Founder/Finance	Infra + hires
.zip codebase blueprint	Engineering	fresh repo intent
ML tuning JSON registry	Internal ML	for feedback

6. 🎁 Pricing Strategy Recommendation (Which Option A/B/D/E/F Is This?)

Ship Option B (sub + packs) first, but stage Option D as differentiated ML feedback economy as a marketing decision for Launch.

Order of pricing to market:

1. **Subscription (Zero)**
 2. 1 × add-on pack upsell shown after day 10 of swiping
 3. 3 × add-on bundling offered after month 2
 4. Enterprise family plan white-label expansion queued at 6-12 months
 5. Sponsored action CPA pricing staged pre-launch with 4.0+ confidence gating
 6. Hybrid reputation model in analytics for pricing + churn elasticity
-

Summary (TL;DR)

- Use **subscription + curated packs** for scalable revenue
 - Increase ARPU with **modular action bundles**
 - License enterprise **white-label** opportunities for stable margin
 - Stage AI tuning + sponsored CPA suggestions for **mid-launch upside**
 - Telemetry captures **cost, confidence, retention, edit rate, fallback usage**
-

If you'd like, tell me and I'll generate:

1. **\$ sizing calculator** for installs & AI call ARPU

2. **Pitch deck** merging models above
3. **10 action sprint breakdown files**
4. **JSON ML registry** with confidence gates
5. **Hiring ramp plan synced to 6–12mo burn**
6. **Investor one-pager and about-us slides**

Which of these should I draft next?

Tab 2

Here's your **full startup-grade business strategy packet** ready for refinement and integration.

1. Investor One-Pager (Concise)

Title: “Zero — Inbox clarity for parents, built with conviction”

Problem

Parents bleed **time and cognitive load** switching tools to decode long email threads and take action on them.

Solution

Zero (iOS) uses:

- AI classification (Mail vs Ads)
- 2–3 sentence thread summaries ≤280 chars
- 1-tap modal actions (top 10 tested for reliability)
- Parent-driven AI tuning loops
- Cost + confidence instrumentation via PostHog

Market

- 100k paying parents → **\$10M+ ARR** potential
- 1M parents → **\$100M+ ARR** category opportunity
- Focus on parents, schools, shoppers, travelers, bill payers

Business Model (Recommended)

- Subscription + packs add-ons + white-label district licensing
- CPA pricing staged for sponsored actions (completion-based)

Why now

Tech lets us build **fast and right** if conviction leads, reducing churn and hallucination risk.

Ask

Seed funding to support:

- Product execution & launch clarity
 - Engineering + AI hires
 - Staged beta (TestFlight) + Marketing install push
 - 30% uncertainty buffer for infra/AI costs
 - 6–12mo runway for proactive IP creation
-

2. Burn + Revenue Coupling (6–12 Month Financial Model)

6-Month Burn (Conservative → Realistic)

Category	6-Month Cost
Eng + Leadership (4 hires)	\$365K–\$550K
QA + Accessibility	\$60K–\$75K
AI tuning (Gemini + classifier + summaries)	\$70K–\$140K
Infrastructure (Cloud Run + Cloud Build)	\$24K–\$66K
Marketing beta + install push	\$270K–\$500K

Legal + compliance + asset creation	\$25K–\$50K
-------------------------------------	-------------

Total 6-mo Burn	\$830K–\$1.37M
------------------------	-----------------------

12-Month Ramp Burn (Conservative Estimate)

Category	12-Month Cost
Engineering + AI hires	\$2.2M–\$2.8M
AI API + Tuning	\$140K–\$280K
Infra	\$66K–\$130K
Research + Retro loops	\$80K–\$120K
Marketing	\$350K–\$550K
Total 12-mo	\$2.5M–\$3.2M

Notes:

- We'll tighten estimates after internal/Top-100 beta telemetry
- We'll add **1 mod suggested action per card** to avoid churn
- AI suggestion bypass is never allowed
- Budget includes 20% buffer + 12.5% “unknown unknowns” hedge

3. 🔄 Staged Beta Plan (iCal-anchored Sprint + Release Order)

1. Internal Dogfood + Cre registry gates (weeks 1–2)

- Reader + summarizer hallucination guardrails
- Top 10 action reliability tests
- Cost & latency instrumentation active

2. Power Parent Beta (weeks 3–6)

- Validate model weights & action override UX

3. TestFlight Rollout (weeks 7–10)

Cohort	Size	Purpose
Alpha dogfood	0–100	Fix missing rights/privacy fields, cost registry
Parents 500	500	Real usage, override AI + summary feedback
Parents 2k	2,000	Pre-launch confidence + performance + accessibility
10k → 50k	10K–50K	Marketing push + passive index stability

4. Launch Hardening + PMF shipping (weeks 11–16)

- 48h stability gate

- Release approval
 - App icon, age/policies, primary category, pricing tiers locked
-

4. Sprint Grooming Plan to Support 10 Core Actions + Summary Quality

Story Point Estimate per action modal set (per 1-week sprint)

Task	Points
Each action modal reliability	8 SP
Classification quality gate	5 SP
Summary hallucination audit	8 SP
AI fallback resiliency	5 SP
Cost + latency registry	3 SP
Accessibility shipping	5 SP
Total/sprint	~39 SP + 20% buffer

5. AI Tuning Prompt Registry for Launch (Copy-Paste Ready for Vertex AI)

Summarizer prompt (version 1.0)

None

You are a concise, accurate email summarizer for busy parents.

Given an email card, return:

1. summary ≤280 chars (2-3 sentences)
2. confidence score (1.0 to 5.0)
3. rationale for confidence
4. 1 best suggested action from 10 core ActionRegistry modals
5. fallback: the best first 2 steps user can take if confidence <3.0 or API fails

Rules:

- Do NOT fabricate details.
- Surface ambiguity explicitly.
- Never send content without human override.

Action inference prompt (version 1.2)

None

You are an action inference engine.

For every email card:

- Pick 1 best action from ActionRegistry.swift "top 10 parent actions"
- Validate requiredContextKeys

Return JSON:

```
{  
  "actionId": string,  
  "displayName": string,  
  "rationale": string,  
  "confidence": number  
}
```

Rules:

- Only 1 action per card.

- Confidence <3.0 → return fallback = “2 most useful first steps”.
- Log conflicts for model tuning (but do NOT store personal attributes).

6. Marketing Campaign + Influencer + Technical Distribution Strategy

Positioning

“Parent life moves fast. Inbox bruises time. Zero gives it back, with AI conviction.”

Campaign Pillars

Pillar	Channel	Output
Install push	App Store + influencer	15–25s reels
Mental model adoption	email, microsite	Beta invite copy
Tech credibility	engineering blog modules	Code cleanliness + docs

Cohort Distribution

- 0–100: Internal team + teens + district prototypes
- 500: Parents + early buyers + co-watchers
- 2k: Pre-launch broader beta
- 10k–50k: Managed marketing rollout

- 1M: Long term white-label partnerships

Tools to Power It

- Install CTA analytics via PostHog
 - Spring UI motion built in SwiftUI
 - Signature capture via PencilKit
 - Pricing/tiers via StoreKit 2
 - Microsite prototypes via React + Tailwind CSS on staging
 - CI/CD via Google Cloud Build and Cloud Run
-

2 Optional Add-on Business Model Slides (To include in 12-Slide Deck)

Slide 1 —

Retention → ARPU unlock model

- Base: \$99.99 annual subscription
- Unlock 1-add-on pack shown after 10 days of swiping

Slide 2 —

Completion Based CPA

- Sponsored action pricing:

- \$2 – \$7 per completed swipe-to-act
 - Only 1 action suggestion visible per card with confidence gating
-

Department Heads for Next Artifact Phase

Department	Ownership	Next Step
Engineering	architecture + modal owner onboarding	Create GitHub repo blueprint
AI/ML	classifier + summarizer tuning loops	Add feedback taxonomy & cost registry
Design	glass UI + convincing .pptx prototypes	Polish 12-Slide storytelling + screens
Marketing	influencer reels + beta invites	Launch install push calculator

If you want next, I can also generate:

1. **Install acquisition calculator (CPC → CPA) in a spreadsheet**
2. **Mock Test Suite prompts for Claude or Vertex AI autoscripts**
3. **Equity hire ramp + SAFE planning section**
4. **Marketing pitch deck + divider frames**
5. **Runway burn planning slide notes**

Say the word and I'll draft all 5 next.

Tab 3

Product Strategy & Business Packet for Zero (6–12mo Runway)

1. Product Requirements Document (Expanded)

Product Vision

Zero turns-email triage into confident, one-swipe actions and accurate summaries for busy parents.

MVP Scope

- **Email pipeline:** Gmail → in-memory processing, no storage.
- **Core AI:** classify (Mail/Ads) + summarize threads/emails.
- **UX:** swipeable cards + 10 high-value modal actions.
- **HIL (Human-in-Loop):** all AI outputs fully editable before execution.
- **Instrumentation:** cost, latency, tokens, confidence.

Success Metrics

Metric	Target
Summary accuracy & relevance	≥ 92%
Classification accuracy	≥ 90%

Action reliability	$\geq 98\%$
AI edit rate	< 30% avg
Latency	<2s local / <3s p95 cloud
Post-launch stability	48h zero critical errors
Time saved	25–40 min/day/user
Beta satisfaction	$\geq 4.3/5$

Target 10 Actions (Registry-Locked for Beta)

- Calendar scheduling / RSVP (EventKit)
- Signature capture (PencilKit)
- Flight check-in flows
- Package tracking
- Invoice payment
- Reminders task creation
- Contact card creation
- List intelligence extraction
- Compound action placeholder resilience
- ML feedback submission

Core Risks + Mitigations

Risk	Impact	Mitigation
Hallucinated summary details	trust erosion	conservative prompts, explicit ambiguity surfacing
Beta crashes	churn	staged TestFlight rollout + crash telemetry
Runway underestimation	hiring stall	bake 6–12mo runway in pricing & ask
Wrong build despite speed	pivot cost	“fast but wrong” guardrails + conviction gating

2. Gantt + 6-Month Staged Beta Plan

6-Month Milestones

None
Month 1–2: Summary + classifier quality gate
Month 3: 500 user managed beta
Month 4: Cost + reliability guardrails finalized
Month 5: 10k marketing soft push
Month 6: Launch + 48h stability gate

Cohort Staging

Stage	Users	Target	Timing
Internal dogfood	0–100	classification+summary reliability	Weeks 1–2

Early parent beta	100–500	HIL override & model agreement	Weeks 3–6
Pre-launch beta	500–2,000	cost+confidence telemetry stabilized	Weeks 7–10
Marketing-led rollout	2,000→10k	broader UX validation	Weeks 11–16
Scale cohort	10k→50k	passive + compound action stability	Weeks 17–24

3. AI Prompt Registry (v1.x series)

Thread Summarization Prompt (v1.1)

None

You are a concise, accurate parent-facing email summarizer.

Return:

1. Summary 2–3 sentences, ≤280 characters
2. Confidence score 1.0–5.0
3. 1 best suggested action from 10-action registry
4. Explicit ambiguity when context incomplete
5. Fallback = 2 most useful first steps if confidence <3.0 or on API error

Rules:

- Do NOT fabricate details.
- User edit + override always stays in control.

Action Inference Prompt (v1.2)

None

You are an action inference engine for email triage.

```
For each email card:  
- Suggest 1 best action from registry  
- Validate required context keys before routing  
- Attach rationale + confidence
```

```
Output JSON only:
```

```
{  
  "actionId": string,  
  "displayName": string,  
  "rationale": string,  
  "confidence": number  
}
```

```
If confidence < 3.0 → include fallbackSteps: [step1, step2]
```

Tuning + Telemetry Requirements

Instrument

Capture

Token usage tokens/request, rolling avg

Performance latency p95, error %, fallback use rate

Cost cost/interaction, CPA estimates

ML agreement AI vs parent overrides → % agreement

Satisfaction modal completion sentiment →
5-score

4. 💰 Runway + Hiring Ramp Cost Model

6-Month Burn (Typical)

Line Item	Cost
1 × Eng Lead	\$180–220K
2 × Senior Swift ICs	\$300–360K
1 × AI/ML Full Stack	\$160–200K
1 × QA/Accessibility IC	\$120–150K
Cloud + CI/CD	\$24–66K
UXR agency for closed beta	\$40–80K
Influencer + Install Push	\$270–500K
Legal/Compliance + GLP subtitles, MSA learnings injection	\$25–50K
Total 6-mo burn	\$1.28M–\$2.18M (keep 20% buffer inside models)

12-Month Burn (Ramp)

Line Item	Cost
Full team ramp + AI tuning + marketing scale	\$2.4M–\$3.3M

5. Business Model Recommendation (3 credible paths)

Path A —

Annual Subscription + Action Packs

- Price: \$99.99/yr + packs (School \$35/yr, Shopping \$45/yr, Travel \$35/yr, Bills \$55/yr)
- Features 1–5 gated into subscription
- Telemetry sold as “Clarity Score” analytics, not extra UX friction
- Strong for retention + modular upsell

Path B —

Free Tier + Pay-As-You-Swipe Completion Credits

- Free base app
- User buys completion credits via StoreKit 2
 - 50 credits = \$14.99
 - 200 credits = \$49.99
 - 1,000 credits = \$199.99
- Credits power the 10 modal actions, plus AI summary confidence improvements
- Good for viral adoption, lower initial friction

Path C —

White-Label District or Busy-Parent OS Licensing

- iOS client stays owned by you, backend licensed to parent ecosystems
- Pricing: \$4/user/month minimum 12 months
- Includes feature 1–5 + reliability SLAs
- Good for long-term scale once metrics validated

Recommended primary go-to-market: Path A, stage Path B and C post-launch to scale quickly without wrong build risk.

6. 🧠 12-Slide Pitch Deck (Narrative + Device Credibility)

Speaker Notes Outline (non-exhaustive, tight 20% conviction narrative)

1. **Cover:** Zero — Inbox clarity for parents
2. **Why Now:** inbox bruises time, AI can fix it if HIL leads
3. **Market:** parents at scale → category \$100M+
4. **Core AI:** 2–3 sentence summaries ≤280 chars
5. **Classifier:** Mail vs Ads ≥90%
6. **Registry:** 10 magical actions, 98% reliability
7. **HIL:** parent override always wins
8. **Telemetry:** cost, latency, confidence
9. **Beta:** dogfood→2k managed rollout
10. **GTM:** influencer + packs upsell
11. **Hiring:** 4 core hires + 20% capacity buffer

12. **Launch:** stability 48h gate, scale next 6–12mo

7. Acceptance Criteria Template (copy-paste)

None

```
Given I am a busy parent
When I open an email card
Then I get a 2–3 sentence summary ≤280 chars
And I get 1 best suggested action from registry
And I can fully edit or override before execution
And latency is <2s local or <3s cloud p95
And the app never stores email data
```

8. Mock QA Claude Prompt (for test suite)

None

```
Audit the Zero iOS codebase and return:
1. Any remaining dead code or TODOs
2. SwiftLint formatting, consistency and schema pass
3. Security audit for auth, keys, and data handling
4. Accessibility pass for VoiceOver + Dynamic Type
5. Error handling and fallback path coverage
6. Test coverage plan ≥80% for top 10 actions
7. Monthly burn and token telemetry wiring checks
```

9. Action Registry CSV (user editable sample)

action	priority
--------	----------

pay_invoice high

rsvp_calendar high

track_package medium
e

(Extend to top 10 in Notion for grooming later.)

Want me to render the deck + spreadsheet fully now?

I can generate polished .docx, .xlsx, .pptx, or .mdk versions next — confirm formats for each artifact:

- PRD?
- Deck?
- Burn Model?
- QA Plan?
- Notion/CSV?

Reply with formats, I'll crank them out now.

Tab 4

Got it — here is your **Zero plan document** in structured text you can paste into a .docx (I kept it crisp and skeletal so it's easy to move to your file).

You can include these directly in the document:

Zero iOS — Next Steps MBA-Anchored Plan

Built with the rigor of a startup-proven product + engineering leader mindset from Harvard MBA.

1. Features Quality Gate (Weeks 1–2)

Goals

- LLM summaries are concise, relevant, and truthful
- Mail vs Ads classifier confidence $\geq 90\%$
- 1-tap actions operate without crashes

Methods

Task	Outcome
Summarization quality — test 250 sample threads	Summaries 2–3 sentences ≤ 280 chars
Classification agreement study (AI vs human swipe)	$\geq 70\%$ agreement baseline
Instrument tokens, latency, cost via PostHog	Full observability per AI call

2. Staged Beta Rollout (TestFlight) (Weeks 3–10)

Tool: use TestFlight platform via App Store Connect.

Stage	Users	Focus
Internal dogfood	0–100	summary, classifier, top 10 actions
Early Parent Beta	100–500	HIL override + model agreement
Pre-launch Beta	500–2,000	cost + latency + completion reliability
	0	
Soft Launch Cohort	2k→10k	scaled UX + passive indexing stability

Required validation

- No AI suggestion friction more than 1 suggestion per card
- Human override always stays in control
- Fallback when confidence < 3.0 or API fails

3. Engineering Hiring Ramp (6–12 Month Horizon)

Roles to hire

- 1 × Engineering Lead (iOS Swift + backend architecture)
- 2 × Senior SwiftUI engineers (modal owners for top 10 actions)
- 1 × AI/ML Full-Stack engineer (classifier + summary tuning + fallback)

- 1 × QA/Accessibility IC (unit/integration tests + accessibility + stability)

Timing

- Eng Lead hired by end of Week 4
 - Swift IC hires Week 6–10
 - QA hire Week 8–12
 - AI hire Week 3–6 (part of infra coupling for beta)
-

4. Marketing + Launch Storytelling (Weeks 2–14)

Campaign pillars

- “Inbox bruises time. Zero gives it back, with AI conviction.”
- 10 magical actions, 1-tap modals, tested by parents
- Invite beta users via email + microsite + influencers

Channels

- Social influencers 15–25s reels
- Parent beta email invites
- App Store launch page subtitle + subtitle
- Microsite hosted via a modern frontend (React + Tailwind CSS) to visually demo flows

CTA scripts

Channel	Script
---------	--------

Install push "Swipe left to act. Tap once. No chaos."

Influencer reel "I swipe once. I see summary. I tap once. Done."

Beta invite email "Your feedback tunes the AI — and you keep control."

KPIs

- Install CPA \$2.50–\$6.00
 - Beta opens $\geq 40\%$
 - Summary edit rate $< 28\%$
 - Reliability 48h gate at launch
-

5. AI Tuning + Model Telemetry Foundation (Parallel to Beta)

Classifier model

Use Gemini 1.5 Flash for classification.

Summarizer model

Use Gemini 1.5 Pro or GPT-4 for summarization.

Prompt registry principles

- 1 best suggestion per card to reduce confusion
- Summaries ≤ 280 characters

- Confidence 1.0–5.0
- Surface ambiguity explicitly, never fabricate
- Provide fallback plan (2 most useful first steps) on confidence <3.0 or API error
- Log conflicts for tuning (no personal or durable attributes included in logs)

Tuning signals to capture

Signal	How it informs tuning
Parent overrides	classifier + action agreement
Summary edits	relevance + hallucination correction
Modal completion sentiment	satisfaction + friction flags
Cost telemetry	pricing + churn elasticity

6. Budget Model (Injecting 1–5 into hiring + marketing + AI burn)

6-month burn estimate: \$1.28M–\$2.18M

12-month ramp burn: \$2.4M–\$3.3M

Includes telemetry buffers for uncertainty, AI API spend, marketing install push, and hiring ramp.

7. Definition of Launch Readiness

None

Given product used by internal team
When 10 actions complete with $\geq 98\%$ reliability
And classification confidence $\geq 90\%$
And summaries achieve $\geq 92\%$ relevance & truthfulness
And cost telemetry is stable and measured
And accessibility audit passes VoiceOver + Dynamic Type
And 48h stability check yields no critical errors
Then the app is ready for broad cohort soft launch.

If you'd like, I can also convert this into a real .docx file for you next — just say "generate docx" and I'll ship it.